

Dr. Thejus Mahajan
Reventlowstraße 17
22605 Hamburg
+49-15259697629
thejus.mahajan@uni-hamburg.de

Hamburg, [PLATZHALTER — Datum der Absendung]

Helmholtz-Zentrum Hereon GmbH
Institut für Umweltchemie des Küstenraumes
Max-Planck-Straße 1
21502 Geesthacht

Application as Postdoctoral Researcher — Probabilistic Deep Learning for Urban Air Quality (AEON-UP) — Reference Code 1056

Dear Dr. Ramacher,

I am writing to apply for the postdoctoral position on probabilistic deep learning for urban air quality. What draws me to AEON-UP is the clause most advertisements leave out: the predictions are to come *with uncertainty estimates*. That question is not an abstract interest for me.

Since finishing my post-doc I have conducted independent research in neural network interpretability, extracting the internal representations of a 15-layer transformer in PyTorch. In the course of it I found two systematic errors in my own pipeline. Both produced entirely plausible output — smooth, well-shaped results that raised no exception and failed no test. One had already reached a published write-up before I caught it, and I corrected it publicly: thejusmahajan.github.io/blog-lc0-attention-frame.html.

The lesson is the one AEON-UP is built around. A confidently wrong model is more dangerous than a visibly uncertain one, and where the output may inform decisions about what people breathe, a calibrated statement of what the model does *not* know is not a refinement — it is the product. Sparse, irregular monitoring networks are exactly the setting where a model must interpolate and declare how far to trust the interpolation.

I would come to this from the physical side as much as the learned one. During my post-doc at Universität Hamburg I developed a plankton life cycle model within the ERGOM framework and ran hindcast and projection simulations for warming scenarios, working daily with gridded spatio-temporal output in NetCDF on Linux HPC systems. I therefore know what a physics-based transport model is, what its output looks like at scale, and where its biases tend to live — which matters when the task is to couple such a model with a learned one rather than replace it. I followed your group's chemistry transport modelling with CMAQ over the North Sea and European domains with real interest, and contributing on the learned side of that coupling is exactly what I want.

I should be straightforward: my experience with Bayesian deep learning and neural processes is at the level of reading and implementation rather than published work. I am building it deliberately — in September I begin the certified course “Deployable Data Analysis & AI Pipelines with HPC” at the Supercomputing-Akademie of the HLRS, following earlier HPC training at the Jülich Supercomputing Centre. What I bring meanwhile is ten years of building computational models and interrogating whether their output can be trusted — and a demonstrated willingness to find that the answer is no.

Hereon is not unfamiliar to me: I worked at the centre as a guest scientist in 2025 on ecosystem modelling, and would be glad to return. I hold work authorisation for Germany, live in Hamburg, and am available for the October start. I would welcome the chance to discuss the project.

Sincerely,



Dr. Thejus Mahajan